



SKILLPILLS

Mini course: Julia
Lecture 1: Introduction

Friederike, Gaston & Juan

material taken/based on Ankur Dhar & James Schloss

friederike.metz@oist.jp

March 1, 2021



- 1 Introduction
 - Why does Julia exist
 - What is wrong with the status quo
- 2 Getting started
- 3 Language syntax and semantics
- 4 The ecosystem

Key Points

- 1 Low-level languages are fast, but high level languages are readable.
- 2 Recent advances in compiler design could bridge these two aspects.
- 3 Julia is being openly developed by researchers, for researchers.

The old mission statement is available [here](#) and some good discussion is available at [here](#).

What kind of language is Julia then?

Julia is a general-purpose, high-level, high-performance, dynamic programming language. It supports parametric polymorphism with multiple dispatch at its core. By default Julia compiles to machine code using "just-ahead-of-time" (JAOT) compilation.

The typical languages used in science are

- 1 Python
- 2 Matlab
- 3 R

Once a problem is becoming too big we usually move to


- 1 C/C++
- 2 Fortran

This is called the 2+ language problem and Julia is trying to solve that.

- The compilers that exist (Numba) only work on primitive types and not user defined ones.
- GIL (Global Interpreter Lock) makes multi-threading hard.
- For fast code you need to write it in C.
- Numpy is great, but awful syntax for math.

- It costs a lot of money and is not open-source.
- Matlab will only be fast for a subset of operations.
- Matlab tends to hide the computer from the programmer.

A (biased) performance comparison



Julia

- is a dynamic language (Julia is easy to use) while having the advantages of static languages (Julia is fast)
- solves the two language problem
- is open-source
- supports parallel and distributed code execution
- is written in Julia (no dependencies to other languages)
- can call other languages
- is made for researchers (but it can do much more)
- comes with many more smart and useful design choices

The Read-Eval-Print-Loop

The REPL is a command-line interface to Julia and is ideal for short experiments.

```

      _          _ _(_) _      | Documentation: https://docs.julialang.org
    (_)          |  (_)  (_)   |
      _ _        | | _  _ _ _   | Type "?" for help, "]"? for Pkg help.
    | | | | | | | | / _ ' |    |
    | | | _ | | | | ( _ | |    | Version 1.5.3 (2020-11-09)
  _ / | \ _ ' _ | | | \ _ _ ' _ | Official https://julialang.org/ release
 | _ _ /                      |

```

julia>

In the REPL you can use ? to switch your REPL mode into help mode and get information about functions.

There are two main IDEs that are *feature* complete and can be used for Julia. The main one is based on **Atom** and is called **Juno**
<http://junolab.org/>.

The second one is based on **Visual Studio Code** and available at
<https://marketplace.visualstudio.com/items?itemName=julialang.language-julia>.

We will not be using these during the course, but you are welcome to try them out.

Jupyter is an interactive web-based client for Python, Julia, R and many other languages. It offers a programming environment that is well suited for explorative data analysis or prototyping.

Installation

```
julia> ENV["JUPYTER"] = ""  
julia> ] add IJulia
```

Starting a Jupyter session

```
julia> using IJulia  
julia> notebook()
```

Other notebook engines

[Pluto.jl](#)

Google Colab ([example template](#))

Julia is a dynamic language and so you can create variables in any scope.

```
x = 1    # x will be of type Int64
y = 1.0  # y will be of type Float64
z = 1.0 - 2.0im # z will be an Complex{Float64}
1//2 # Rational numbers
"This is a String"
"""
This is a multiline
String
"""
'C' # Character literal
1.0f0 # Float32 literal
#=
This is a multiline comment
=#
```

Use `typeof(x)` to check the type of any variable `x`.

Variable names can be unicode and so greek symbols can be used. In the REPL and most editors you can insert them by entering their \LaTeX name and press [Tab].

Exercise

Create a variable λ that stores the value of π . Check the type of λ . Look up the docs for the `convert` function. Convert λ to float.

Variable names can be unicode and so greek symbols can be used. In the REPL and most editors you can insert them by entering their \LaTeX name and press [Tab].

Exercise

Create a variable λ that stores the value of π . Check the type of λ . Look up the docs for the `convert` function. Convert λ to float.

Solution

```
# \lambda + <tab> = \pi + <tab>
lambda = pi
typeof(lambda)
?convert
convert(Float32, lambda)
```

```
# Tuple: (item1, item2, ...)
okinawa = ("Onna", "Nago", "Naha")
okinawa[1]

# NamedTuple: (name1 = item1, name2 = item2, ...)
okinawa = (central = "Onna", north = "Nago", south = "Naha")
okinawa.north

# Dictionaries: Dict{key1 => value1, key2 => value2, ...}
password = Dict{"Jeremie" => "12345", "Juan" => nothing}
password["Juan"] = "4444"

# Arrays: [item1, item2, ...]
fibonacci = [1, 1, "two", 3, "five", 8, 13]
fibonacci[3] = 2
```

Julia is 1-based indexing, not 0-based like Python!!!
You can add or remove elements using `push!` and `pop!`.

Julia has all the typical conditionals `if`, `else`, `elseif` which have to end in an `end`. Blocks in Julia are not whitespace sensitive and conditionals do not need to be wrapped in round brackets.

```
if rand() < 0.5
    println("Gaston is awesome!")
elseif rand() >= 0.5
    println("Jeremie is awesome!")
else
    println("Juan is awesome!")
end

# Ternary operator: a ? b : c
rand() < 0.5 ? println("Gaston is awesome!") : println("Juan
    is awesome!")
```

`&&` is the and operator, `||` is the or operator.

Exercise

Print an error message if a variable x is greater than zero

Exercise

Print an error message if a variable `x` is greater than zero

Solution

```
if (x > 0) error("x cannot be greater than 0") end

# Ternary operator
x > 0 ? error("x cannot be greater than 0") : nothing

# Short-circuit evaluation
(x > 0) && error("x cannot be greater than 0")
```

Julia has for and while loops. A while loop takes a condition and a for loop takes an iterator. One can use break to break out of a loop and continue to skip to the next iteration.

```
for (i, x) in enumerate(['A', 'B', 'C'])
    if x == 'B'
        continue
    end
    println(i)
end

while true
    rand() < 0.1 ? break : println("You are trapped!")
end
```

Exercise

Create a dictionary that holds integers (up to the value 100) and their squares as key, value pairs.

Exercise

Create a dictionary that holds integers (up to the value 100) and their squares as key, value pairs.

Solution

```
squares = Dict()
for i in 1:100
    squares[i] = i^2
end
println(squares)

# Array comprehension
squares_array = [i^2 for i in 1:100]
```

Julia uses functions to organize operations. Every function is compiled for the combination of input parameters.

```
function f(x, y)
    return x + y
end

g(x) = x^2
h = x -> 1/x

# map takes a function as input and applies that function to
# every element of the data structure you pass it.
map(lowercase, ['A', 'B', 'C'])
map(x->x+2, [1, 2, 3])
```

By convention, functions followed by ! alter their contents and functions lacking ! do not, e.g. `sort()` vs. `sort!()`.

The `::` operator annotates the type of variables. For example, `x::Int8` means that the value of `x` is an instance of `Int8`. There are abstract types (`Integer`, `Number`) and concrete types (`Int8`, `Float64`) and their relationships form a type graph.

In Julia a function is a set of multiple methods. Each method is defined for particular argument types and comes with its own implementation. When you call a function the most specific method is executed (dispatched).

```
h(x::Number) = println("x is most definitely a number.")  
  
h(x::Integer) = println("x is an integer")  
  
h(x::Int8) = println("Specific method for Int8")
```

Multiple dispatch makes your code generic and fast!

You can create your own types using struct.

```
abstract type Coordinate end

mutable struct Cartesian <: Coordinate
    position::Tuple{Float64, Float64}
end

struct Polar <: Coordinate
    radius::Float64
    phi::Real
end

point = Polar(1,pi)
point.radius
```

Keep in mind: Julia is not an Object-Oriented programming language so functions do not belong to an object!

Exercise

Create a struct `Point` with two field names `x`, `y` both being of type `Real`. Implement a function `add` that takes two arguments of type `Point` and returns the sum of their `x`'s and the difference of their `y`'s.

Exercise

Create a struct `Point` with two field names `x`, `y` both being of type `Real`. Implement a function `add` that takes two arguments of type `Point` and returns the sum of their `x`'s and the difference of their `y`'s.

Solution

```
struct Point
    x::Real
    y::Real
end

function add(arg1::Point, arg2::Point)
    return arg1.x + arg2.x, arg1.y - arg2.y
end
```

```
# Specifying a return type will convert the returned value
g(x, y)::Int8 = x * y

# Varargs functions (variable number of arguments functions)
h(a,b,x...) = (a,b,x)

# Optional arguments
f(x, y=1, z::Int64=2) = x + y + z

# Keyword arguments (seperated by a semicolon)
g(x, y; a=2, b::String="goya") = x+y+a, b
```

Everything above can be combined arbitrarily, e.g. keyword arguments can be collected using ...

There are mainly two different ways to run your julia scripts:

```
# in your command prompt or terminal
```

```
julia script.jl
```

```
# inside the REPL
```

```
include("script.jl")
```

Working from the REPL is preferred since all methods inside the script are compiled the first time and stay available as long as the session is not closed.

Try to break your programs down into lots of small specialized functions to really take advantage of the julia compiler and its speed-up.

If your julia files are becoming too large, organize them into modules (namespaces). Module names are capitalised and you can also nest them.

```
module MyModule
    export f

    g() = "Internal function"
    f() = println(g())
end

# inside the REPL
include("MyModule.jl")

# Loading the module in a different file
using Main.MyModule
```

Additional advantage: Modules can be precompiled and saved to hard-drive which means precompiled code won't disappear when closing the REPL.

Julia has an inbuilt package manager called Pkg. You can access it from the REPL by entering `]`. For example, to install and load the Julia package `Distributions.jl` in the REPL just run:

```
# Installing a package
] add Distributions
# Loading it afterwards
using Distributions
```

A few other commands:

```
# Updating all installed packages
] update
# What packages are installed?
] status
```

Look [here](#) for more information about Julia packages.

Documentation <https://docs.julialang.org/en/v1.5/>

Forum <https://discourse.julialang.org>

GitHub <https://github.com/JuliaLang/julia>

Slack <https://julialang.org/slack/>

YouTube <https://www.youtube.com/user/JuliaLanguage>

Learning <https://julialang.org/learning/>

OIST also has a Slack workspace: `julia@oist`

Question?!

What do you want to hear/learn more about?

Wednesday How to use Julia for computation and plotting

Friday Problem sets and ???

Pseudocode

- ❶ For each of N iterations
 - ❶ Select a random point inside a square of area $4r^2$ as Cartesian, (x, y) , coordinates. (You can choose $r = 1$ for simplicity.)
 - ❷ Determine if the point also falls inside the circle embedded within this square of area πr^2 .
 - ❸ Keep track of whether or not this point fell inside the circle. At the end of N iterations, you want to know M – the number of the N random points that fell inside the circle!
- ❷ Calculate π as $4\frac{M}{N}$

Solution

```
function calculate_pi(N::Integer = 1000)
    M = 0
    for i in 1:N
        x = 1-2*rand()
        y = 1-2*rand()
        # Check if the distance to (x, y) from (0, 0) is less than
        # the radius of 1
        if x^2 + y^2 < 1
            M += 1
        end
    end
    return 4 * M / N
end

map(calculate_pi, [10, 100, 1000, 10000, 100000])
```