



Mini course: Julia

Lecture 3: IDEs, Data Structures, Coding Practices, and Algorithms

Gaston, Friederike,
and Juan

Okinawa Institute of Science
and Technology

March 5, 2021



1 IDEs & remote workflow

```
if condition == true  
|   f()  
end
```



2 Data Structures

```
if condition  
|   f()  
end
```



3 Coding Practices

```
if(condition) f() end
```



4 Algorithms

```
var = ifelse(condition,f(),nothing)
```



```
condition ? f() : nothing
```



```
condition && f()
```



When designing this Skill Pill, we assumed the following

- ▶ You have your own workflow in regards to which software and libraries you work with.
- ▶ You have seen and are familiar with common data structures.
- ▶ You know how to program and use programming as part of your daily work.

As such, we have designed today's lesson so that you may begin using Julia in your work as soon as possible.

Workflow



Julia in VS Code



- ▶ Step 1: Install Julia.
- ▶ Step 2: Run]add IJulia on Julia's REPL.
- ▶ Step 3: Install Anaconda distribution
- ▶ Step 4: Launch Jupyter lab/notebook from the Anaconda Navigator
- ▶ Step 5: Create a new jupyter notebook with Julia kernel.

- ▶ Step 1: Install Julia, and Atom.
- ▶ Step 2: Search in Atom for the packages: `uber-juno`, `language-julia`, `julia-client`, `ink`, and `indent-detective`.
- ▶ Step 3: Restart and launch Julia in the REPL inside Atom. It will gather all remaining necessary packages.

- ▶ Step 1: Install Julia, and Visual Studio Code.
- ▶ Step 2: Inside VSCode, go to extensions and search "julia". Install Julia for VSCode and restart.

Depending on your workflow, you may want to have access to data, more computational resources, or just to avoid ~~stressing out the crappy~~ Macbooks using your personal laptop for running research code.

On Jupyter notebook, this is done very easily. Make sure you have installed Jupyter notebook on both the local and remote computers.

Assuming you have access (via VPN) and know the IP of your workstation to run the following ...

On remote host terminal:

```
$ jupyter notebook --no-browser --port=8889
```

In the **local host**, run the following in the terminal (cmd for Windows):

To connect:

```
$ ssh -N -f -L localhost:8888:localhost:8889 user@remote_host_IP
```

Make sure to change *user* to your real username in **remote host**, and change *your_remote_host_name* to your address of your working station.

Example:

```
$ ssh -N -f -L localhost:8888:localhost:8889 gaston@10.2.80.145
```

Now open web browser and type:

Web browser address bar:

localhost:8888

In Atom, the setup for remote workflow is very simple.

- ▶ Install the following package inside Atom: **ftp-remote-edit**.
- ▶ Go to Packages → Ftp-Remote-Edit → Toggle.
- ▶ Right click on Remote tab → Edit server.
- ▶ Name the server and input the necessary details.
- ▶ To use tmux. Go to **julia-client** package settings → Remote settings → Use persistent tmux session. Also check that the command to execute Julia remotely is appropriate.
- ▶ Start the remote server by going to Juno → Start Remote Julia Process.

Remote connection to your workstation

Ftp-Remote-Edit Server Settings

You can edit each connection at the time. All changes will only be saved by pushing the save button.

server@office	New	Delete	Duplicate
The name of the server.	Test		
server@office	- None -	Edit	
The hostname or IP address of the server.	Port		
10.2.80.102	22		
Protocol			
SFTP - SSH File Transfer Protocol			
Logon Type			
Username / Password			
Username for authentication.			
gsivori			
Password/Passphrase for authentication.			

Initial Directory.			
/home/gsivori/Dropbox (OIST)/Documents/PhD-Thesis/Project1			
<input type="button" value="Save"/>	<input type="button" value="Import"/>	<input type="button" value="Cancel"/>	

DataStructures.jl has the following data structures:

- ▶ Deque (based on block-list)
- ▶ Stacks and Queues
- ▶ Accumulators and Counters
- ▶ Disjoint Sets
- ▶ Binary Heap
- ▶ Mutable Binary Heap
- ▶ Ordered Dicts and Sets
- ▶ Dictionaries with Defaults
- ▶ Trie (Tree)
- ▶ Linked List
- ▶ Sorted Dict, Multi-Dict and Set
- ▶ Priority Queue

DataStructures.jl has the following data structures:

- ▶ Deque (based on block-list)
- ▶ Stacks and Queues
- ▶ Accumulators and Counters
- ▶ Disjoint Sets
- ▶ Binary Heap
- ▶ Mutable Binary Heap
- ▶ Ordered Dicts and Sets
- ▶ Dictionaries with Defaults
- ▶ Trie (Tree)
- ▶ Linked List
- ▶ Sorted Dict, Multi-Dict and Set
- ▶ Priority Queue

All information regarding these Data Structures can be found Here:
<http://datastructuresjl.readthedocs.io/en/latest/index.html/>

All of these algorithms can be viewed with @edit. We'll use two more before hopping into algorithms:
Binary Trees and Priority Queues

Based on contiguous blocks on memory. Similar to Julia's **Vector** implementation, but, as it grows, it avoids copying all values to a new Vector.

```
a = Deque{Int}()  
isempty(a)           # test whether the dequeue is empty  
length(a)           # get the number of elements  
push!(a, 10)         # add an element to the back  
pop!(a)              # remove an element from the back  
pushfirst!(a, 20)    # add an element to the front  
popfirst!(a)         # remove an element from the front  
first(a)             # get the element at the front  
last(a)              # get the element at the back
```

Circular Buffer

A circular buffer of fixed capacity. New items are pushed to the back of the list, overwriting values circularly.

```
cb = CircularBuffer{Int}(n) # n-element circular buffer
isfull(cb) # test whether the buffer is full
isempty(cb) # test whether the buffer is empty
empty!(cb) # reset the buffer
capacity(cb) # return capacity
size(cb) # same as length(cb)
push!(cb, 10) # overwrites front if full
pop!(cb) # remove the element at the back
pushfirst!(cb, 10) # overwrites back if full
popfirst!(cb) # remove the element at the front
append!(cb, [1, 2, 3, 4]) # push at most last `capacity` items
convert(Vector{Float64}, cb) # convert items to type Float64
eltype(cb) # return type of items
fill!(cb, data) # grows the buffer up-to capacity
```

```
s = Stack{Int}()      # create a stack
eltype(s)            # get the type of elements
push!(s, 1)          # push back a item
first(s)             # get an item from the top of stack
pop!(s)              # get and remove a first item
iterate(s::Stack)   # Get a LIFO iterator of a stack
Iterators.reverse(s::Stack{T}) # Get a FILO iterator of a stack
s1 == s2 # check whether the two stacks are same
```

```
q = Queue{Int}()
enqueue!(q, x)
x = first(q)
x = last(q)
x = dequeue!(q)
```

PriorityQueue



A type of data structure that enqueues items by priority.

```
julia> # Julia code
        pq = PriorityQueue();
julia> # Insert keys with associated priorities
        pq["a"] = 10; pq["b"] = 5; pq["c"] = 15; pq
PriorityQueue{Any,Any,Base.Order.ForwardOrdering} with 3 entries:
  "b" => 5
  "a" => 10
  "c" => 15
julia> # Change the priority of an existing key
        pq["a"] = 0; pq
PriorityQueue{Any,Any,Base.Order.ForwardOrdering} with 3 entries:
  "a" => 0
  "b" => 5
  "c" => 15
```

A type of self-balancing binary tree that has an extra bit of information to make searching, inserting, and deleting operations with $O(\log(n))$:
https://en.wikipedia.org/wiki/Red%20black_tree

```
julia> tree = RBTree{Int}();
julia> for k in 1:2:20
        push!(tree, k)
    end
julia> haskey(tree, 3)
true
julia> tree[4]
7
julia> for k in 1:2:10
        delete!(tree, k)
    end
julia> haskey(tree, 5)
false
```

Some recommended coding practices to consider with Julia.

- ▶ Measure performance with `@time` and check for memory allocations.
- ▶ Break a function into multiple definitions (proper use of multiple dispatcher)
- ▶ Consider variable type stability for compiler optimizations.
- ▶ Access arrays in memory order (along columns in Julia)
- ▶ Pre-allocate output.
- ▶ Use vectorized operation "`@.`" for readability.
- ▶ Consider using `@views` to look at big arrays.
- ▶ Read through `Profile.jl`, `BenchmarkingTools.jl`, and `Cthulu.jl` packages for further optimizations and debugging.
- ▶ Differences with other languages: <https://docs.julialang.org/en/v1/manual/noteworthy-differences/>

@time

Measure performance with @time. Notice how using global variable allocates unnecessary memory making this example run longer.

```
julia> x = rand(1000);
julia> function sum_global()
           s = 0.0
           for i in x
               s += i
           end
           return s
       end;
julia> @time sum_global()
 0.017705 seconds (15.28 k allocations: 694.484 KiB)
496.84883432553846
julia> @time sum_global()
 0.000140 seconds (3.49 k allocations: 70.313 KiB)
496.84883432553846
```

Compared to passing it as an argument:

```
julia> x = rand(1000);
julia> function sum_arg(x)
           s = 0.0
           for i in x
               s += i
           end
           return s
       end;
julia> @time sum_arg(x)
 0.007701 seconds (821 allocations: 43.059 KiB)
496.84883432553846
julia> @time sum_arg(x)
 0.000006 seconds (5 allocations: 176 bytes)
496.84883432553846
```

```
using LinearAlgebra
function mynorm(A)
    if isa(A, Vector)
        return sqrt(real(dot(A,A)))
    elseif isa(A, Matrix)
        return maximum(svdvals(A))
    else
        error("mynorm: invalid argument")
    end
end
```

This is better:

```
norm(x::Vector) = sqrt(real(dot(x, x)))
norm(A::Matrix) = maximum(svdvals(A))
```

Consider variable type

```
function foo()
    x = 1
    for i = 1:10
        x /= rand()
    end
    return x
end
```

```
function foo()
    x::Float64 = 1.0 / rand()
    for i = 2:10
        x /= rand()
    end
    return x
end
```

Multi-dimensional arrays are stored in column-major order (arrays are stacked one column at a time).

```
julia> x = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4
julia> x[:, :]
4-element Array{Int64,1}:
 1
 3
 2
 4
```

A rule of thumb to keep in mind is that with column-major arrays, the first index changes most rapidly. Looping will be faster if the inner-most loop index is the first to appear in a slice expression.

Pre-allocate output

If your function returns an Array or some other complex type, it may have to allocate memory. Unfortunately, often times allocation and its converse, garbage collection, are substantial bottlenecks. Compare this:

```
julia> function xinc(x)
           return [x, x+1, x+2]
       end;
julia> function loopinc()
           y = 0
           for i = 1:10^7
               ret = xinc(i)
               y += ret[2]
           end
           return y
       end;
```

Pre-allocate output

To this:

```
julia> function xinc!(ret::AbstractVector{T}, x::T) where T
    ret[1] = x
    ret[2] = x+1
    ret[3] = x+2
    nothing
end;
julia> function loopinc_prealloc()
    ret = Vector{Int}(undef, 3)
    y = 0
    for i = 1:10^7
        xinc!(ret, i)
        y += ret[2]
    end
    return y
end;
```

Pre-allocate output

This changes a lot.

```
julia> @time loopinc()
0.529894 seconds (40.00 M allocations: 1.490 GiB, 12.14% gc
    time)
50000015000000
julia> @time loopinc_prealloc()
0.030850 seconds (6 allocations: 288 bytes)
50000015000000
```

Use vectorized operation for readability



Using . for broadcast opearation messes up readability. Use "@." for vectorized operations.

```
julia> f(x) = 3x.^2 + 4x + 7x.^3;
julia> fdot(x) = @. 3x^2 + 4x + 7x^3
```

If x is an Array, fdot(x) will have significantly faster performance.

```
julia> x = rand(10^6);
```

```
julia> fcopy(x) = sum(x[2:end-1]);
```

```
julia> @views fview(x) = sum(x[2:end-1]);
```

```
julia> @time fcopy(x);
0.003051 seconds (7 allocations: 7.630 MB)
```

```
julia> @time fview(x);
0.001020 seconds (6 allocations: 224 bytes)
```

Exercise

Implement your favorite algorithm in Julia

I will be coding the Lorentz Attractor as an example. Lorentz function given by:

$$\begin{aligned}\frac{dx}{dt} &= \sigma(y - x) \\ \frac{dy}{dt} &= x(\rho - z) - y \\ \frac{dz}{dt} &= xy - \beta z\end{aligned}$$

Solve this for $t = 1 : 150$. Initial conditions $[1, 1, 1]$.
 $\sigma = 10, \rho = 28, \beta = 8/3$.

The end of this mini course!



Thanks for joining us!